

Reinforcement Learning

Yasuyuki Sawada, Yaolang Zhong

University of Tokyo

`yaolang.zhong@e.u-tokyo.ac.jp`

Recap: Motivation for Machine Learning

- ▶ Dynamic programming faces the curse of dimensionality
 - ▶ Parameterization problem
 - ▶ Classical functional forms (linear, polynomials, splines)
 - ▶ Number of basis terms grows rapidly with state dimension
 - ▶ Strong restrictions and poor capture of interactions
 - ▶ Grid / data problem
 - ▶ Grid-based DP requires N^d state points
 - ▶ Computation and memory grow exponentially
 - ▶ Sparse grids mitigate but do not eliminate the problem

How Machine Learning Addresses the Curse of Dimensionality

- ▶ Machine learning relaxes both constraints
 - ▶ Neural networks for parameterization
 - ▶ Flexible, high-dimensional function approximation
 - ▶ Captures nonlinearities and interactions
 - ▶ Stochastic simulation for data collection
 - ▶ Learn from simulated trajectories, not full grids
 - ▶ Scales to high-dimensional state spaces

Overview

1. Network Architecture

- ▶ Actor: policy network $\sigma(s; \theta)$ maps states to actions
- ▶ Critic: value network $Q(s, a; \phi)$ evaluates state–action pairs

2. Forward Pass

- ▶ Batch of states $S \in \mathbb{R}^{N \times D_S}$
- ▶ Actor produces actions; critic produces value estimates

3. Objective Functions

- ▶ Actor: maximize expected return via policy gradient
- ▶ Critic: minimize Bellman residual (temporal difference error)

4. Backward Pass and Parameter Updates

- ▶ Compute gradients $\nabla_{\theta} \mathcal{L}_{\theta}$ and $\nabla_{\phi} \mathcal{L}_{\phi}$
- ▶ Update parameters via gradient descent


5. Implementation Caveats

- ▶ Target networks for training stability
- ▶ Gradient detachment to separate actor and critic updates


Actor–Critic Methods: Introduction

- ▶ Actor–critic methods are a general framework for reinforcement learning (RL) problem
- ▶ Particularly suited for Large or continuous state spaces and action spaces
- ▶ The framework decomposes learning into policy learning (actor) and value learning (critic)
 - ▶ Actor: a parameterized policy $\sigma(a \mid s; \theta)$ that selects actions
 - ▶ Critic: a parameterized value function $Q(s, a; \phi)$ that provides an evaluation signal guiding policy improvement
 - ▶ The notation $\sigma(a \mid s; \theta)$ denotes a (possibly stochastic) policy: the probability of choosing action a in state s given parameters θ ; deterministic policies are a special case that assign probability one to a single action (often written simply as $\sigma(s)$)

Actor Network Architecture



Critic Network Architecture



Forward Pass

- ▶ Suppose a batch of states available: $\{s_i\}_{i=1}^N$ where $s_i \in \mathbb{R}^{D_S}$, stacked as $S \in \mathbb{R}^{N \times D_S}$
- ▶ Two neural networks:
 - ▶ Actor: $\sigma(s; \theta)$
 - ▶ Critic: $Q(s, a; \phi)$
- ▶ Forward Pass:

$$\begin{aligned} S \in \mathbb{R}^{N \times D_S} &\xrightarrow{\theta} A \in \mathbb{R}^{N \times D_A} \\ (S, A) \in \mathbb{R}^{N \times (D_S + D_A)} &\xrightarrow{\phi} Q \in \mathbb{R}^N \end{aligned}$$

- ▶ Training vs. evaluation:
 - ▶ In *training mode*, the forward pass constructs a computational graph and records gradients with respect to the network parameters
 - ▶ In *evaluation mode*, the same forward mappings are applied, but no gradients are recorded

Backward Propagation: Intuition

- ▶ In training mode, once a scalar loss $\mathcal{L}$ is defined, the gradient $\nabla_{\theta}\mathcal{L}$ is computed automatically by back-propagation in any ML package.
- ▶ Backpropagation applies the chain rule through the computational graph constructed in the forward pass
- ▶ The actor and critic would be updated via gradient descent (although in practice, more stochastic methods are used) as:

$$\theta \leftarrow \theta - \alpha_{\theta} \nabla_{\theta} \mathcal{L}_{\theta}$$

$$\phi \leftarrow \phi - \alpha_{\phi} \nabla_{\phi} \mathcal{L}_{\phi}$$

- ▶ The question is how to construct the loss function $\mathcal{L}$ for the actor and critic.

Backward Propagation: Actor Objective (1/2)

- ▶ The actor aims to maximize expected discounted returns:

$$J(\theta) = \mathbb{E}_{\tau \sim P_{\theta}}[R(\tau)], \quad R(\tau) \equiv \sum_{k=0}^{\infty} \gamma^k r_k$$

- ▶ The policy gradient can be written as:

$$\nabla_{\theta} J(\theta) = \mathbb{E} \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \sigma(a_t | s_t; \theta) \gamma^t G_t \right], \quad G_t \equiv \sum_{k=t}^{\infty} \gamma^{k-t} r_k$$

- ▶ Realized returns G_t are noisy and depend on the entire future trajectory
- ▶ For a given policy $\sigma(a | s; \theta)$, the state-action value satisfies:

$$Q^{\sigma}(s_t, a_t) = \mathbb{E}[G_t | s_t, a_t]$$

Backward Propagation: Actor Objective (2/2)

- ▶ By iterated expectations, replacing G_t with $Q^\sigma(s_t, a_t)$ is exact in expectation:

$$\mathbb{E}[\nabla_\theta \log \sigma(a_t | s_t; \theta) G_t] = \mathbb{E}[\nabla_\theta \log \sigma(a_t | s_t; \theta) Q^\sigma(s_t, a_t)]$$

- ▶ In practice, the true Q^σ is unknown; we use the current critic estimate $Q(s, a; \phi)$
- ▶ Sample-based actor loss (depends on current critic parameters ϕ):

$$\mathcal{L}_\theta = -\frac{1}{N} \sum_{i=1}^N \sum_{t=1}^{T_i} \log \sigma(a_{i,t} | s_{i,t}; \theta) \gamma^t Q(s_{i,t}, a_{i,t}; \phi)$$

- ▶ The critic $Q(s, a; \phi)$ provides a lower-variance learning signal compared to realized returns

Policy Gradient Derivation* (1/4: Likelihood-Ratio Reformulation)

- ▶ Start from the RL objective in integral form:

$$J(\theta) = \mathbb{E}_{\tau \sim P_\theta}[R(\tau)] = \int R(\tau) P_\theta(\tau) d\tau, \quad \text{where} \quad R(\tau) \equiv \sum_{k=0}^{\infty} \gamma^k r_k$$

- ▶ Differentiate under the integral sign:

$$\nabla_\theta J(\theta) = \int R(\tau) \nabla_\theta P_\theta(\tau) d\tau$$

- ▶ Use $\nabla_\theta P_\theta(\tau) = P_\theta(\tau) \nabla_\theta \log P_\theta(\tau)$:

$$\nabla_\theta J(\theta) = \int R(\tau) P_\theta(\tau) \nabla_\theta \log P_\theta(\tau) d\tau$$

- ▶ Convert to expectation:

$$\nabla_\theta J(\theta) = \mathbb{E}_{\tau \sim P_\theta} [\nabla_\theta \log P_\theta(\tau) R(\tau)]$$

This is the likelihood-ratio (score-function) trick.

Policy Gradient Derivation* (2/4: Trajectory Likelihood Factorization)

- ▶ Trajectory density factorizes as

$$P_{\theta}(\tau) = p(s_0) \prod_{t=0}^{\infty} \sigma(a_t \mid s_t; \theta) P(s_{t+1} \mid s_t, a_t)$$

- ▶ Only the policy depends on θ , hence

$$\nabla_{\theta} \log P_{\theta}(\tau) = \sum_{t=0}^{\infty} \nabla_{\theta} \log \sigma(a_t \mid s_t; \theta)$$

Environment transition terms drop out of the gradient.

Policy Gradient Derivation* (3/4: Causality)

- ▶ Substitute the log-policy sum and expand $R(\tau)$:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim P_{\theta}} \left[\left(\sum_{t=0}^{\infty} \nabla_{\theta} \log \sigma(a_t | s_t; \theta) \right) \left(\sum_{k=0}^{\infty} \gamma^k r_k \right) \right]$$

- ▶ Swap the order of summation:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim P_{\theta}} \left[\sum_{t=0}^{\infty} \sum_{k=0}^{\infty} \nabla_{\theta} \log \sigma(a_t | s_t; \theta) \gamma^k r_k \right]$$

- ▶ **Causality:** For each t , terms with $k < t$ have zero expectation, since rewards before t do not depend on a_t . Therefore, restrict inner sum to $k \geq t$:

$$= \mathbb{E} \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \sigma(a_t | s_t; \theta) \sum_{k=t}^{\infty} \gamma^k r_k \right]$$

- ▶ Define the return from t : $G_t \equiv \sum_{k=t}^{\infty} \gamma^{k-t} r_k$, so $\sum_{k=t}^{\infty} \gamma^k r_k = \gamma^t G_t$

Policy Gradient Derivation* (4/4: Surrogate Objective)

- ▶ Substitute G_t and collect terms:

$$\nabla_{\theta} J(\theta) = \mathbb{E} \left[\sum_{t=0}^{\infty} \nabla_{\theta} \log \sigma(a_t | s_t; \theta) \gamma^t G_t \right]$$

- ▶ **Surrogate objective:**

$$\max_{\theta} \mathbb{E} \left[\sum_{t=0}^{\infty} \log \sigma(a_t | s_t; \theta) \gamma^t G_t \right]$$

- ▶ This surrogate is constructed so that its gradient with respect to θ equals the true policy gradient $\nabla_{\theta} J(\theta)$.
- ▶ We maximize the surrogate because it is differentiable from sampled data and enables gradient-based optimization of the policy.
- ▶ Sample-based approximation:

$$\hat{J}(\theta) = \frac{1}{N} \sum_{i=1}^N \sum_{t=1}^{T_i} \log \sigma(a_{i,t} | s_{i,t}; \theta) \gamma^t G_{i,t}$$

Backward Propagation: Critic Objective

- ▶ The critic aims to approximate the state–action value function $Q^\sigma(s, a)$
- ▶ The true Q^σ satisfies the Bellman equation:

$$Q^\sigma(s_t, a_t) = r_t + \gamma \mathbb{E}_{s_{t+1}}[Q^\sigma(s_{t+1}, a_{t+1})]$$

- ▶ Sample-based critic loss (mean squared Bellman error):

$$\mathcal{L}_\phi = \frac{1}{N} \sum_{i=1}^N (Q(s_i, a_i; \phi) - (r_i + \gamma \mathbb{E}_{s_{i+1}}[Q(s_{i+1}, a_{i+1}; \phi)]))^2$$

Implementation Caveat 1: Target Networks

- ▶ Problem: The critic loss depends on its own predictions

$$\mathcal{L}_\phi = \frac{1}{N} \sum_{i=1}^N \left(Q(s_i, a_i; \phi) - \underbrace{(r_i + \gamma Q(s_{i+1}, a_{i+1}; \phi))}_{\text{target uses same } \phi} \right)^2$$

- ▶ This creates a *moving target* problem $\Rightarrow$ training instability
- ▶ Solution: Use separate *target networks* $\phi_{\text{target}}, \theta_{\text{target}}$

$$y_i = r_i + \gamma Q(s_{i+1}, \sigma(s_{i+1}; \theta_{\text{target}}); \phi_{\text{target}})$$

- ▶ Target networks are updated slowly:
 - ▶ Hard update: Copy parameters every K steps
 - ▶ Soft update (Polyak averaging): $\phi_{\text{target}} \leftarrow \tau \phi + (1 - \tau) \phi_{\text{target}}$, with $\tau \ll 1$
- ▶ Result: Stable targets $\Rightarrow$ more stable learning


Implementation Caveat 2: Detaching the Action

The `detach` operation creates a copy of a tensor that is disconnected from the computational graph, blocking gradient propagation.

- ▶ The action $a = \sigma(s; \theta)$ is produced by the actor
- ▶ During critic updates, the action should be treated as fixed input data
- ▶ Detaching prevents gradients from flowing back to θ

Implementation:

- ▶ `a_detached = a.detach()`
- ▶ Pass `a_detached` as input to $Q(s, a; \phi)$



The critic update optimizes ϕ only; the actor is updated separately via policy gradient.