

Lecture 8: Simulation-Based Learning

Yasuyuki Sawada, Yaolang Zhong

University of Tokyo
yaolang.zhong@e.u-tokyo.ac.jp

December 10, 2025

Recap: The Computational Challenge


- ▶ Recall the Bellman Equation:

$$V(s) = \max_a \left\{ r(s, a) + \beta \int V(s') P(ds' | s, a) \right\}$$

- ▶ So far, deterministic update:
 - ▶ enumerate all states $s \in \mathcal{S}$ (use a fixed grid in the continuous case)
 - ▶ update $V(s)$ by enumerating over all states $s' \in \mathcal{S}$ and integrating over P
 - ▶ Lecture 7: introduce stochasticity by Monte Carlo integration (but still iterate over all s)
- ▶ However: Enumerating all states s means we want accurate $V(s)$ everywhere: the evaluation and update is costly, but the contribution to gain in accuracy is marginal.
- ▶ Ergodic set: the set of states that the agent visits infinitely often under a particular policy. Denote the ergodic set corresponding to a policy π as $\bar{\mathcal{S}}_\pi$.

Example 1: Ergodic Set in a Discrete Markov Chain


Consider a 9-state MDP. The optimal policy π^* induces transitions that only recurrently visit states $\{1, 2, 3, 4\}$.



- ▶ Even for a suboptimal policy π , once it improves sufficiently, we often have $\bar{\mathcal{S}}_\pi = \bar{\mathcal{S}}_{\pi^*}$.
- ▶ Implication: updating $V(s)$ for $s \notin \bar{\mathcal{S}}_{\pi^*}$ wastes computation—these states are never revisited.


Example 2: Ergodic Set in a Growth Model

In the stochastic growth model $(k_t, z_t) \rightarrow (k_{t+1}, z_{t+1})$, the ergodic set is a small region of the state space.



Judd, Maliar, Maliar & Valero (2014): Simulated trajectories visit $< 5\%$ of grid points in high-dimensional models. Updating only visited states achieves similar accuracy at a fraction of the cost.

Planning vs Simulation: illustration



- ▶ Trajectory / Episode / Rollout: A sequence of states and rewards $(s_0, a_0, r_0, s_1, a_1, r_1, \dots)$ generated by following a policy π from initial state s_0 to terminal state s_T .
- ▶ The associated return is $G_0 = r_0 + \beta r_1 + \beta^2 r_2 + \dots + \beta^{T-1} r_{T-1}$

Planning: The Bellman Update

- At iteration k , we update $V^{(k)}$ to $V^{(k+1)}$ using the Bellman equation:

$$V^{(k+1)}(s) = r(s) + \beta \int V^{(k)}(s') P(ds'|s) = \underbrace{(r(s) + \beta V^{(k)}(s'))}_{\text{Target}} P(ds'|s)$$

- Planning computes the integral exactly (quadrature) or via MC integration with M samples:

$$V^{(k+1)}(s) = \frac{1}{M} \sum_{m=1}^M [r(s) + \beta V^{(k)}(s'_m)] = \frac{1}{M} \sum_{m=1}^M \text{Target}_m^{(k)}$$

- Simulation-based: sample size $M = 1$ and incremental update weighted by α_k (learning rate):

$$\begin{aligned} V^{(k+1)}(s) &= V^{(k)}(s) + \alpha_k \underbrace{[r + \beta V^{(k)}(s') - V^{(k)}(s)]}_{\text{Bellman residual } \delta_k} \\ &= (1 - \alpha_k) \cdot \underbrace{V^{(k)}(s)}_{\text{old estimate}} + \alpha_k \cdot \underbrace{[r(s) + \beta V^{(k)}(s')]}_{\text{Target}} \end{aligned}$$

Convergence: Simulation → Planning

- ▶ Suppose state s has been visited $n(s)$ times. On the $n(s)$ -th visit, choose $\alpha_{n(s)} = \frac{1}{n(s)}$, which satisfies the Robbins–Monro conditions: $\sum_k \alpha_k = \infty$ and $\sum_k \alpha_k^2 < \infty$.
- ▶ The simulation update on the $n(s)$ -th visit becomes

$$V^{(k+1)}(s) = V^{(k)}(s) + \frac{1}{n(s)} \left[\underbrace{r(s) + \beta V^{(k)}(s') - V^{(k)}(s)}_{\text{Target}_{n(s)}} \right] = \frac{1}{n(s)} \text{Target}_{n(s)} + \frac{n(s)-1}{n(s)} V^{(k)}(s)$$

- ▶ Unrolling: substitute $V^{(k)}(s) = \frac{1}{n(s)-1} \text{Target}_{n(s)-1} + \frac{n(s)-2}{n(s)-1} V^{(k-1)}(s)$, and repeat:

$$\begin{aligned} V^{(k+1)}(s) &= \frac{1}{n(s)} \text{Target}_{n(s)} + \frac{n(s)-1}{n(s)} \cdot \frac{1}{n(s)-1} \text{Target}_{n(s)-1} + \dots \\ &= \frac{1}{n(s)} \text{Target}_{n(s)} + \frac{1}{n(s)} \text{Target}_{n(s)-1} + \dots + \frac{1}{n(s)} \text{Target}_1 \\ &= \frac{1}{n(s)} \sum_{i=1}^{n(s)} \text{Target}_i \end{aligned}$$

Temporal Difference (TD) Learning: Formal Introduction

- ▶ The simulation-based update we derived uses the estimated value $V(s_{t+1})$ to update $V(s_t)$. This is called *bootstrapping*: updating an estimate based on another estimate. (This differs from the bootstrap in statistics, which is resampling from data to estimate uncertainty.)
- ▶ Simulation-based methods with bootstrapping are called Temporal Difference (TD) learning:

$$V(s_t) \leftarrow V(s_t) + \alpha \cdot (r_t + \beta V(s_{t+1}) - V(s_t)), \quad \text{where } s_t \sim \bar{\mathcal{S}}_\pi$$

- ▶ The name “Temporal Difference” refers to the difference between value estimates at two successive time points: $V(s_{t+1})$ vs $V(s_t)$.
- ▶ The Bellman residual $\delta_t = r_t + \beta V(s_{t+1}) - V(s_t)$ is also called the TD error.
- ▶ This specific update rule is called TD(0). The “0” indicates that we bootstrap immediately after observing just one reward r_t , using zero additional steps before estimating the remaining value via $V(s_{t+1})$.

TD(0) Algorithm for Policy Evaluation

- ▶ Step 0 (inputs): Policy π to evaluate; discount $\beta \in (0, 1)$; learning rate schedule $\{\alpha_k\}$; number of episodes K ; initial state distribution μ_0 .
- ▶ Step 1 (initialize): For all states $s \in \mathcal{S}$, set $V^{(0)}(s) \leftarrow 0$ (or arbitrary values).
- ▶ Step 2 (episode loop): For episode $k = 1, 2, \dots, K$:
 - ▶ Sample initial state $s_0 \sim \mu_0$.
 - ▶ For $t = 0, 1, 2, \dots$ until terminal state:
 1. Take action $a_t = \pi(s_t)$, observe reward r_t and next state s_{t+1} .
 2. Compute TD error: $\delta_t = r_t + \beta V(s_{t+1}) - V(s_t)$.
 3. Update: $V(s_t) \leftarrow V(s_t) + \alpha_k \cdot \delta_t$.
- ▶ Step 3 (output): Value function estimate $V \approx V^\pi$.

The Exploration Problem: Ergodic Set Depends on π

- ▶ The ergodic set of states visited under a policy,

$$\bar{\mathcal{S}}_{\pi} = \{s : \Pr_{\pi}(s_t = s \text{ for some } t) > 0\},$$

depends on the policy π itself.

- ▶ During learning, the policy evolves:

$$\pi^{(0)} \rightarrow \pi^{(1)} \rightarrow \pi^{(2)} \rightarrow \dots \rightarrow \pi^*.$$

- ▶ This creates a fundamental difficulty:


- ▶ Accurate estimation of $V(s)$ requires many visits to state s .
- ▶ But the states we visit depend on the current policy $\pi^{(k)}$.
- ▶ A poor early policy may prevent us from ever reaching states that matter under π^* .

- ▶ Consequences of poor initialization:

- ▶ $\bar{\mathcal{S}}_{\pi^{(0)}}$ may omit states in $\bar{\mathcal{S}}_{\pi^*}$.
- ▶ For such states, $V(s)$ is never updated.
- ▶ Policy improvement cannot discover these states, and learning becomes trapped in a suboptimal region.

Example: How Poor Initialization Causes Suboptimal Convergence

Same 9-state MDP, but now consider a poorly initialized policy $\pi^{(0)}$ that gets stuck in $\{5, 6, 9\}$.



- ▶ Policy $\pi^{(0)}$ cycles within $\{5, 6, 9\}$ — never transitions to state 1.
- ▶ $V(s)$ for $s \in \{1, 2, 3, 4\}$ never gets updated \Rightarrow cannot discover higher rewards there.
- ▶ Learning is stuck: $\bar{\mathcal{S}}_{\pi^{(0)}} \cap \bar{\mathcal{S}}_{\pi^*} = \emptyset$.

Solutions to the Exploration Problem

- ▶ ε -greedy policy:
 - ▶ With probability $1 - \varepsilon$: follow the current best policy
 - ▶ With probability ε : take a random action
 - ▶ Ensures all state-action pairs have positive visitation probability
- ▶ Optimistic initialization:
 - ▶ Initialize $V(s)$ to high values (optimistic about unexplored states)
 - ▶ Agent is “curious” — explores until reality disappoints the optimism
- ▶ Exploring starts:
 - ▶ Begin episodes from randomly chosen states
 - ▶ Guarantees coverage of the full state space over time
- ▶ Key insight: Exploration is the price we pay for not knowing the model. Planning (with known P) can evaluate all states without visiting them; simulation cannot.

TD Learning: Advantages and Disadvantages

- ▶ Advantages:
 - ▶ Model-free: No need to know the transition probabilities P .
 - ▶ Scalable: DP must iterate over all states each sweep; TD only updates visited states, making it feasible for large or continuous state spaces.
 - ▶ Computationally cheaper: No integration required
- ▶ Disadvantages :
 - ▶ Noisy updates: using one sample s' instead of integrating over all possible s' introduces variance.
 - ▶ Exploration required: Only visited states get updated.
 - ▶ Slower convergence: Stochastic approximation converges at $O(1/\sqrt{n})$; DP with known model converges geometrically.
 - ▶ Learning rate tuning: Performance depends on α . DP has no such hyperparameter.

Generalizing TD(0): The n -Step Return

- ▶ Recall TD(0) bootstraps after observing just one reward:

$$G_t^{(1)} = r_t + \beta V(s_{t+1})$$

- ▶ What if we wait longer? Observe n actual rewards before bootstrapping:

$$G_t^{(n)} = r_t + \beta r_{t+1} + \beta^2 r_{t+2} + \cdots + \beta^{n-1} r_{t+n-1} + \beta^n V(s_{t+n})$$

- ▶ The n -step return uses n realized rewards and bootstraps from $V(s_{t+n})$.

- ▶ Examples:

- ▶ $n = 1$ (TD(0)): $G_t^{(1)} = r_t + \beta V(s_{t+1})$
- ▶ $n = 2$: $G_t^{(2)} = r_t + \beta r_{t+1} + \beta^2 V(s_{t+2})$
- ▶ $n = 3$: $G_t^{(3)} = r_t + \beta r_{t+1} + \beta^2 r_{t+2} + \beta^3 V(s_{t+3})$

- ▶ As n increases, we rely more on actual rewards and less on the estimated value.

The Two Extremes: TD(0) and Monte Carlo

- ▶ TD(0) ($n = 1$): Bootstrap immediately after one reward.


$$G_t^{(1)} = \underbrace{r_t}_{\text{1 actual reward}} + \underbrace{\beta V(s_{t+1})}_{\text{bootstrap}}$$

- ✓ Low variance: only one random transition affects the target.
- ✗ Biased: relies on current (imperfect) estimate $V(s_{t+1})$.
- ▶ Monte Carlo ($n = T - t$, i.e., wait until episode ends): No bootstrapping at all.

$$G_t^{(\infty)} = \underbrace{r_t + \beta r_{t+1} + \beta^2 r_{t+2} + \cdots + \beta^{T-t} r_T}_{\text{all actual rewards, no bootstrap}}$$

- ✓ Unbiased: $G_t^{(\infty)}$ is a true sample of $V(s_t)$.
- ✗ High variance: depends on many random transitions.
- ✗ Must wait for episode to end before updating.

Visualizing the Spectrum: Backup Diagrams



- ▶ TD(0): Bootstrap after 1 step (red arrow shows bootstrap).
- ▶ n -step: Bootstrap after n steps.
- ▶ MC: No bootstrap—use full trajectory to terminal state s_T .

TD(0) vs MC: The Bias-Variance Trade-off

- ▶ TD(0) is biased and MC is unbiased:
 - ▶ TD(0): $r_t + \beta V(s_{t+1})$ is based on the estimated value $V(s_{t+1})$, which is our current guess—not the true value. We are “updating a guess towards another guess.”
 - ▶ MC: base on the actual realized rewards, which is a random variable whose expectation equals the true $V(s_t)$.
- ▶ TD(0) has lower variance than MC:
 - ▶ TD(0): depends on only one random transition (s_t, r_t, s_{t+1}) .
 - ▶ MC: accumulates randomness from all transitions in the episode:
 $(s_t, r_t, s_{t+1}), (s_{t+1}, r_{t+1}, s_{t+2}), \dots$
- ▶ The n -step return interpolates: larger n means less bias but more variance.

TD(λ): A Weighted Blend of All n -Step Returns

- ▶ Instead of choosing a single n , why not combine all n -step returns?
- ▶ The λ -return takes a weighted average with geometrically decaying weights:

$$G_t^\lambda = (1 - \lambda) \sum_{n=1}^{T-t-1} \lambda^{n-1} G_t^{(n)} + \lambda^{T-t-1} G_t^{(T-t)}$$

- ▶ The parameter $\lambda \in [0, 1]$ controls the decay:
 - ▶ Weight on $G_t^{(1)}$: $(1 - \lambda)$
 - ▶ Weight on $G_t^{(2)}$: $(1 - \lambda)\lambda$
 - ▶ Weight on $G_t^{(3)}$: $(1 - \lambda)\lambda^2$
 - ▶ \vdots
- ▶ Special cases:
 - ▶ $\lambda = 0$: All weight on $G_t^{(1)}$ \Rightarrow TD(0).
 - ▶ $\lambda = 1$: All weight on $G_t^{(\infty)}$ \Rightarrow Monte Carlo.

Numerical Experiment: VFI vs TD Policy Evaluation

Setup: Evaluate the *same* optimal policy π^* (solved by VFI) using different methods.

- ▶ Model: Consumption-saving with discrete assets and income shocks.
- ▶ Ground truth: V^{VFI} from exact dynamic programming.
- ▶ TD evaluation: Run $\text{TD}(\lambda)$ for $\lambda \in \{0, 0.5, 0.9, 1.0\}$.
- ▶ Repeat each TD method 5 times with different random seeds.

Metrics:

- ▶ MSE: $\frac{1}{|\mathcal{S}|} \sum_s (V^{\text{TD}}(s) - V^{\text{VFI}}(s))^2$ (bias + variance)
- ▶ Variance: $\frac{1}{|\mathcal{S}|} \sum_s \text{Var}[V^{\text{TD}}(s)]$ across runs

Key insight: All methods evaluate the *same* policy—differences are purely due to estimation method.

Code: Comparing VFI and TD(λ)

```
from model import create_default_model
from algos import compare_vfi_td

model = create_default_model()
results = compare_vfi_td(
    model,
    lambdas=[0.0, 0.5, 0.9, 1.0],
    n_runs=5,
    n_episodes=5000,
    alpha=0.5
)

# Results contain:
# - V_vfi: exact value function from VFI
# - td_results[lambda]: V_mean, V_std, mse, avg_variance
```

See Lab_8_Simulation_Based_Learning/ for full implementation.

Expected Results: Bias-Variance Trade-off

Method	MSE vs VFI	Avg Variance	Interpretation
TD(0), $\lambda = 0$	Medium	Low	High bias, low variance
TD(0.5)	Lower	Medium	Balanced
TD(0.9)	Lower	Higher	Less bias, more variance
MC, $\lambda = 1$	Lowest bias	Highest	Unbiased but noisy

Observations:

- ▶ TD(0) converges fast but may have residual bias (depends on α schedule).
- ▶ MC is unbiased in the limit but requires many episodes to reduce variance.
- ▶ Intermediate λ (e.g., 0.9) often achieves the best MSE in practice.
- ▶ All TD methods are slower than VFI when the model is known—but TD works without knowing P .