

# Lecture 13: Heterogeneous Agent Model

Yasuyuki Sawada, Yaolang Zhong

University of Tokyo  
yaolang.zhong@e.u-tokyo.ac.jp

January 21, 2026

## Recap: Motivation for Machine Learning

- ▶ Dynamic programming faces the curse of dimensionality
  - ▶ Parameterization problem
    - ▶ Classical functional forms (linear, polynomials, splines)
    - ▶ Number of basis terms grows rapidly with state dimension
    - ▶ Strong restrictions and poor capture of interactions
  - ▶ Grid / data problem
    - ▶ Grid-based DP requires  $N^d$  state points
    - ▶ Computation and memory grow exponentially
    - ▶ Sparse grids mitigate but do not eliminate the problem

# How Machine Learning Addresses the Curse of Dimensionality

- ▶ Machine learning relaxes both constraints
  - ▶ Neural networks for parameterization
    - ▶ Flexible, high-dimensional function approximation
    - ▶ Captures nonlinearities and interactions
  - ▶ Stochastic simulation for data collection
    - ▶ Learn from simulated trajectories, not full grids
    - ▶ Scales to high-dimensional state spaces

## RL for Economic Models: Data Generation

- ▶ One observation:  $(s, a, r, s')$ 
  - ▶  $s$ : taken as given
  - ▶  $a$ : determined by policy  $\sigma_\theta(s)$
  - ▶  $r, s'$ : determined by the environment's reward function  $R(s, a)$  and transition function  $P(s' | s, a)$
- ▶ Policy Evaluation (Update  $Q_\phi^\sigma$ ) — requires  $\sigma_\theta$  from Policy Improvement:
  - ▶ Fixed-point iteration:

$$Q_\phi^\sigma(s, a) \leftarrow R(s, a) + \beta \sum_{s'} P(s' | s, a) Q_\phi^\sigma(s', \sigma_\theta(s'))$$

- ▶ Gradient-based update:

$$\phi \leftarrow \phi - \alpha \nabla_\phi \underbrace{\left( R(s, a) + \beta \sum_{s'} P(s' | s, a) Q_\phi^\sigma(s', \sigma_\theta(s')) - Q_\phi^\sigma(s, a) \right)^2}_{\text{Loss}}$$

## RL for Economic Models: Q-Learning / Value Iteration

- ▶ One observation:  $(s, a, r, s')$
- ▶ Q-Learning (Update  $Q_\phi$ ) — self-contained, only  $Q_\phi$  needed:
  - ▶ Fixed-point iteration:

$$Q_\phi(s, a) \leftarrow R(s, a) + \beta \sum_{s'} P(s' | s, a) \max_{a'} Q_\phi(s', a')$$

- ▶ Gradient-based update:

$$\phi \leftarrow \phi - \alpha \nabla_\phi \underbrace{\left( R(s, a) + \beta \sum_{s'} P(s' | s, a) \max_{a'} Q_\phi(s', a') - Q_\phi(s, a) \right)^2}_{\text{Loss}}$$

## RL for Economic Models: Policy Improvement

- ▶ One observation:  $(s, a, r, s')$
- ▶ Policy Improvement (Update  $\sigma_\theta$ ) — requires  $Q_\phi$  from Policy Evaluation:
  - ▶ Fixed-point iteration:

$$\sigma_\theta(s) \leftarrow \arg \max_a Q_\phi(s, a)$$

- ▶ Gradient-based update:

$$\theta \leftarrow \theta + \alpha \nabla_\theta \underbrace{\log \sigma_\theta(a | s) \cdot A(s, a)}_{\text{Policy Gradient}}$$

where  $A(s, a) = Q_\phi(s, a) - \mathbb{E}_{a \sim \sigma_\theta}[Q_\phi(s, a)]$  is the advantage

## RL for Economic Models: Euler Equation Update

- ▶ One observation:  $(s, a, r, s')$  where  $a = c$  (consumption),  $s = (k, z)$
- ▶ Euler Equation Residual (Update  $c_\theta$ ) — self-contained, only  $c_\theta$  needed:
  - ▶ Fixed-point iteration:

$$u'(c_\theta(s)) = \beta \sum_{s'} P(s' | s, a) [1 + r(s') - \delta] u'(c_\theta(s'))$$

- ▶ Gradient-based update:

$$\theta \leftarrow \theta - \alpha \nabla_\theta \underbrace{\left( u'(c_\theta(s)) - \beta \sum_{s'} P(s' | s) R'(s') u'(c_\theta(s')) \right)^2}_{\text{Euler Residual Loss}}$$

where  $R'(s') = 1 + r(s') - \delta$  is the gross return

## State Sampling Methods

- ▶ Updates require samples  $(s, a, r, s')$ ; given  $s$ , the tuple is determined by  $(\sigma_\theta, R, P)$
- ▶ Three approaches to obtain state  $s$ :
  1. Grid-based (non-stochastic): Deterministic enumeration over  $\mathcal{S}$ 
$$s \in \{s^{(1)}, s^{(2)}, \dots, s^{(N)}\}$$

2. Ergodic sampling: Draw from stationary distribution

$$s \sim \mu^* \quad (\text{long-run distribution under } \sigma^*)$$

3. Trajectory sampling: Sequential states from simulation

$$(s_0, a_0, r_0, s_1, a_1, r_1, s_2, \dots) \quad \text{where } s_{t+1} \sim P(\cdot | s_t, a_t)$$

# State Sampling Methods: Comparison

|              | Grid-Based         | Ergodic          | Trajectory        |
|--------------|--------------------|------------------|-------------------|
| Coverage     | Full $\mathcal{S}$ | Likely states    | Path-dependent    |
| Correlation  | None               | None (i.i.d.)    | High (sequential) |
| Scalability  | $O(N^d)$           | Good             | Good              |
| Requirements | Grid construction  | Estimate $\mu^*$ | Simulation        |

- ▶ Grid-based: Classical DP approach; curse of dimensionality
- ▶ Ergodic: Concentrates on economically relevant states; requires  $\mu^*$  estimation
- ▶ Trajectory: Standard RL approach; requires decorrelation via replay buffer

## Experience Replay Buffer

- ▶ Problem: data generated by sequential interaction is highly correlated
  - ▶ Violates i.i.d. assumptions underlying gradient-based optimization
  - ▶ Leads to unstable and inefficient learning
- ▶ Solution: store past transitions in a replay buffer  $\mathcal{D} = \{(s, a, r, s')\}_{t=1}^N$
- ▶ Training loop:
  1. Interact with the environment using current policy  $\pi_\theta$  (eval mode)
  2. Store observed transitions in  $\mathcal{D}$
  3. Sample a random mini-batch from  $\mathcal{D}$
  4. Update actor and critic using the sampled batch
- ▶ Key benefits:
  - ▶ Breaks temporal correlation and stabilizes training
  - ▶ Improves sample efficiency by reusing past experience

# Online (On-Policy) vs. Offline (Off-Policy) Learning

- ▶ On-policy: Data generated by the current policy  $\pi_\theta$ 
  - ▶ Learning uses state distribution  $\mu^{\pi_\theta}$ ; policy and data evolve jointly
  - ▶ Advantages: Stable learning, no distribution mismatch
  - ▶ Limitations: Data discarded after each update; low sample efficiency
- ▶ Off-policy: Data collected by a different or past policy
  - ▶ Learning target decoupled from data collection
  - ▶ Advantages: High sample efficiency via replay; flexible data reuse
  - ▶ Limitations: Potential instability due to distribution mismatch
- ▶ Implications for economic models:
  - ▶ Simulation is cheap; training neural networks is expensive
  - ▶ Off-policy learning preferred to reuse simulated data efficiently

# Parallel Simulation for Data Collection

- ▶ Vectorized environments: Simulate  $M$  agents in parallel

$$\{s_t^{(1)}, s_t^{(2)}, \dots, s_t^{(M)}\} \xrightarrow{\pi_\theta} \{a_t^{(1)}, a_t^{(2)}, \dots, a_t^{(M)}\}$$

- ▶ Each step generates  $M$  transitions simultaneously
- ▶ Benefits:
  - ▶ GPU parallelism for policy evaluation
  - ▶ Faster data collection
  - ▶ Better gradient estimates (more samples per update)

## Heterogeneous Agent Models

- ▶ Representative agent models assume identical agents
  - ▶ Cannot capture wealth/income inequality
  - ▶ Miss redistributive effects of policy
  - ▶ Aggregate consumption = individual consumption (scaled)
- ▶ Heterogeneous agent models allow:
  - ▶ Idiosyncratic risk: individual-specific shocks (unemployment, health)
  - ▶ Incomplete markets: cannot fully insure against idiosyncratic risk
  - ▶ Endogenous wealth distribution: emerges from optimization
- ▶ Krusell & Smith (1998): First model combining HA with aggregate uncertainty

## Krusell-Smith (1998): Key Innovation

- ▶ Challenge: With aggregate shocks, agents need to forecast future prices
- ▶ Future prices depend on future aggregate capital  $K'$ , which in terms future wealth distribution  $\mu'$
- ▶ Future wealth distribution  $\mu'$  depends on the current wealth distribution  $\mu$  and the policy function  $\sigma$
- ▶ The wealth distribution  $\mu$  is infinite-dimensional
- ▶ Krusell-Smith insight:
  - ▶ Agents use bounded rationality
  - ▶ Approximate  $\mu$  with a few moments (e.g., mean capital  $\bar{K}$ )
  - ▶ Empirically: first moment captures 99.9% of price variation

## Model Environment: Preferences

- ▶ Continuum of agents indexed by  $i \in [0, 1]$
- ▶ Time is discrete, infinite horizon:  $t = 0, 1, 2, \dots$
- ▶ Preferences over consumption streams:

$$\mathbb{E}_0 \left[ \sum_{t=0}^{\infty} \beta^t u(c_{it}) \right]$$

- ▶ Utility function (CRRA):

$$u(c) = \frac{c^{1-\gamma} - 1}{1 - \gamma}, \quad \gamma > 0$$

- ▶  $\beta \in (0, 1)$ : discount factor
- ▶  $\gamma$ : coefficient of relative risk aversion

## Model Environment: Labor Income Process

- ▶ Aggregate state  $z \in \{z_g, z_b\}$  (good/bad times)
  - ▶ Follows Markov chain with transition matrix  $\Pi^z$
  - ▶ Affects aggregate productivity
- ▶ Idiosyncratic state  $\epsilon \in \{0, 1\}$  (unemployed/employed)
  - ▶ Transition probabilities depend on aggregate state  $z$
  - ▶  $\pi(\epsilon' | \epsilon, z, z')$ : probability of  $\epsilon'$  given current  $(\epsilon, z)$  and next  $z'$
- ▶ Labor income:

$$y_{it} = \begin{cases} w_t \bar{\ell} & \text{if } \epsilon_{it} = 1 \text{ (employed)} \\ \mu^u & \text{if } \epsilon_{it} = 0 \text{ (unemployed, UI)} \end{cases}$$

where  $w_t$  is the wage and  $\bar{\ell}$  is labor supply

## Model Environment: Budget Constraint

- ▶ Agents can save in a single asset: capital  $a_{it} \geq 0$
- ▶ Borrowing constraint:  $a_{it} \geq \underline{a}$  (often  $\underline{a} = 0$ )
- ▶ Budget constraint:

$$c_{it} + a_{it+1} = (1 + r_t - \delta)a_{it} + y_{it}$$

where:

- ▶  $r_t$ : rental rate of capital
- ▶  $\delta$ : depreciation rate
- ▶  $y_{it}$ : labor income
- ▶ Incomplete markets: No insurance against  $\epsilon$  shocks
- ▶ Agents self-insure through precautionary savings

## Production Technology

- ▶ Representative firm with Cobb-Douglas technology:

$$Y_t = z_t K_t^\alpha L_t^{1-\alpha}$$

- ▶  $K_t = \int a_{it} di$ : aggregate capital
- ▶  $L_t = \int \epsilon_{it} \bar{\ell} di$ : aggregate effective labor
- ▶ Competitive factor markets imply:

$$r_t = \alpha z_t \left( \frac{K_t}{L_t} \right)^{\alpha-1}$$

$$w_t = (1 - \alpha) z_t \left( \frac{K_t}{L_t} \right)^\alpha$$

## State Variables

- ▶ Individual state:  $(a, \epsilon)$ 
  - ▶  $a$ : individual asset holdings
  - ▶  $\epsilon$ : employment status
- ▶ Aggregate state:  $(z, \mu)$ 
  - ▶  $z$ : aggregate productivity shock
  - ▶  $\mu$ : distribution of agents over  $(a, \epsilon)$
- ▶ Key challenge:  $\mu$  is infinite-dimensional
- ▶ Prices  $(r, w)$  are functions of  $(z, \mu)$

## Household's Recursive Problem

- ▶ Value function  $V(a, \epsilon; z, \mu)$  solves:

$$V(a, \epsilon; z, \mu) = \max_{c, a'} \{ u(c) + \beta \mathbb{E} [V(a', \epsilon'; z', \mu') \mid \epsilon, z] \}$$

- ▶ Subject to:

$$\begin{aligned} c + a' &= (1 + r(z, \mu) - \delta)a + y(\epsilon, w(z, \mu)) \\ a' &\geq \underline{a} \\ \mu' &= \Gamma(z, \mu, z') \end{aligned}$$

- ▶  $\Gamma$ : law of motion for the distribution (unknown!)
- ▶ Policy function:  $a' = g(a, \epsilon; z, \mu)$

## Law of Motion for Distribution

- ▶ Given policy function  $g$  and transition  $\pi$ :

$$\mu'(A, \epsilon') = \int_{a \in A} \int_{\epsilon} \mathbf{1}\{g(a, \epsilon; z, \mu) \in A\} \cdot \pi(\epsilon' | \epsilon, z, z') d\mu(a, \epsilon)$$

- ▶ This is a functional equation in  $\mu$
- ▶ Aggregate capital evolves:

$$K' = \int a' d\mu' = \int g(a, \epsilon; z, \mu) d\mu(a, \epsilon)$$

## Recursive Competitive Equilibrium: Definition

A recursive competitive equilibrium consists of:

1. Value function  $V(a, \epsilon; z, \mu)$
2. Policy function  $g(a, \epsilon; z, \mu)$
3. Price functions  $r(z, \mu), w(z, \mu)$
4. Law of motion  $\Gamma(z, \mu, z')$

Such that:

- (i) Given prices and  $\Gamma$ ,  $V$  and  $g$  solve household's problem
- (ii) Prices satisfy firm's FOCs:

$$r = \alpha z \left( \frac{K}{L} \right)^{\alpha-1}, \quad w = (1 - \alpha)z \left( \frac{K}{L} \right)^{\alpha}$$

- (iii) Markets clear:  $K = \int a d\mu$
- (iv)  $\Gamma$  is consistent with  $g$  and transition probabilities

## The Curse of Dimensionality

- ▶ Problem: Distribution  $\mu$  is infinite-dimensional
- ▶ Cannot store  $\mu$  on a computer directly
- ▶ Standard approaches fail:
  - ▶ Discretize  $a$  into  $N_a$  grid points
  - ▶ With  $N_\epsilon = 2$  employment states
  - ▶ Need to track  $N_a \times N_\epsilon$  probabilities
  - ▶ Value function becomes  $V(a, \epsilon; z, \mu_1, \dots, \mu_{N_a \times N_\epsilon})$
- ▶ Even with  $N_a = 100$ : state space is  $\mathbb{R}^{200}$
- ▶ Infeasible to solve by standard dynamic programming

## Krusell-Smith's Bounded Rationality Approach

- ▶ Key insight: Agents don't know the true law of motion  $\Gamma$
- ▶ Instead, agents use a perceived law of motion (PLM)
- ▶ Approximate  $\mu$  with moments:  $\vec{m} = (m_1, m_2, \dots, m_J)$
- ▶ Simplest case:  $\vec{m} = \bar{K}$  (mean capital only)
- ▶ PLM for log aggregate capital:

$$\log K' = \begin{cases} a_0 + a_1 \log K & \text{if } z = z_g \\ b_0 + b_1 \log K & \text{if } z = z_b \end{cases}$$

- ▶ Coefficients  $(a_0, a_1, b_0, b_1)$  to be determined in equilibrium

# Approximate Equilibrium

An approximate recursive competitive equilibrium consists of:

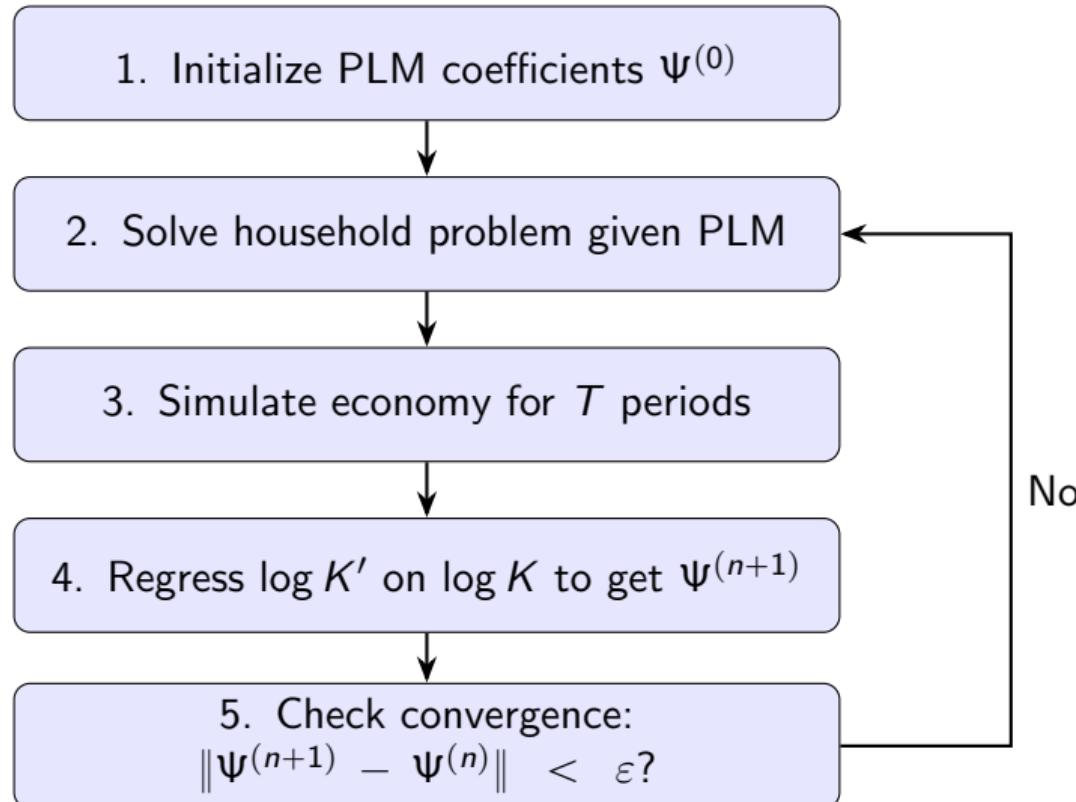
1. Value function  $\tilde{V}(a, \epsilon; z, K)$
2. Policy function  $\tilde{g}(a, \epsilon; z, K)$
3. Price functions  $\tilde{r}(z, K), \tilde{w}(z, K)$
4. PLM coefficients  $\Psi = (a_0, a_1, b_0, b_1)$

Such that:

- (i) Given prices and PLM,  $\tilde{V}$  and  $\tilde{g}$  solve household's problem
- (ii) Prices from firm's FOCs with  $K$  as capital stock
- (iii) PLM  $\Psi$  provides a good forecast of actual  $K'$

Accuracy criterion:  $R^2 > 0.9999$  for forecasting  $K'$

## Krusell-Smith Algorithm: Overview



## Step 1: Initialize PLM

- ▶ Set initial guess for PLM coefficients:

$$\Psi^{(0)} = (a_0^{(0)}, a_1^{(0)}, b_0^{(0)}, b_1^{(0)})$$

- ▶ Common starting point:
  - ▶  $a_1^{(0)} = b_1^{(0)} \approx 0.99$  (high persistence)
  - ▶  $a_0^{(0)}, b_0^{(0)}$  chosen so steady-state  $K$  is consistent
- ▶ Or use steady-state Aiyagari model as initial guess

## Step 2: Solve Household Problem

Given PLM  $\Psi$ , solve the Bellman equation:

$$\tilde{V}(a, \epsilon; z, K) = \max_{a' \geq a} \left\{ u(c) + \beta \sum_{z', \epsilon'} \pi(z', \epsilon' | z, \epsilon) \tilde{V}(a', \epsilon'; z', K') \right\}$$

where:

$$\begin{aligned} c &= (1 + r(z, K) - \delta)a + y(\epsilon, w(z, K)) - a' \\ K' &= H(z, K; \Psi) \quad (\text{from PLM}) \end{aligned}$$

Solution methods:

- ▶ Value function iteration on grid
- ▶ Endogenous grid method (faster)
- ▶ Policy function iteration

## Step 3: Simulate the Economy

- ▶ Track  $I$  agents for  $T$  periods (e.g.,  $I = 10,000$ ,  $T = 11,000$ )
- ▶ Initialize: draw  $(a_0^i, \epsilon_0^i)$  from ergodic distribution
- ▶ For each period  $t$ :
  1. Compute aggregate capital:  $K_t = \frac{1}{I} \sum_{i=1}^I a_t^i$
  2. Draw aggregate shock  $z_{t+1}$  from Markov chain
  3. For each agent  $i$ :
    - ▶ Draw  $\epsilon_{t+1}^i$  from  $\pi(\cdot | \epsilon_t^i, z_t, z_{t+1})$
    - ▶ Compute  $a_{t+1}^i = \tilde{g}(a_t^i, \epsilon_t^i; z_t, K_t)$
- ▶ Discard first  $T_0$  periods (burn-in, e.g.,  $T_0 = 1,000$ )

## Step 4: Update PLM via Regression

- ▶ From simulation, obtain time series  $\{K_t, z_t\}_{t=T_0}^T$
- ▶ Run OLS regressions separately for each  $z$ :

$$\log K_{t+1} = a_0 + a_1 \log K_t + \varepsilon_t \quad \text{for } t : z_t = z_g$$

$$\log K_{t+1} = b_0 + b_1 \log K_t + \varepsilon_t \quad \text{for } t : z_t = z_b$$

- ▶ New PLM coefficients:  $\Psi^{(n+1)} = (\hat{a}_0, \hat{a}_1, \hat{b}_0, \hat{b}_1)$
- ▶ Check  $R^2$  for forecast quality (should be  $> 0.9999$ )

## Step 5: Check Convergence

- ▶ Convergence criterion:

$$\|\Psi^{(n+1)} - \Psi^{(n)}\| < \varepsilon$$

- ▶ Typically  $\varepsilon = 10^{-4}$  or smaller

- ▶ If not converged:

- ▶ Update:  $\Psi^{(n+1)} = \lambda \Psi^{(n+1)} + (1 - \lambda) \Psi^{(n)}$
- ▶ Dampening factor  $\lambda \in (0.3, 0.5)$  for stability
- ▶ Return to Step 2

- ▶ If converged:

- ▶ Verify  $R^2 > 0.9999$  for both regressions
- ▶ This validates the approximate equilibrium

## The Near-Aggregation Property

- ▶ Surprising result: First moment ( $\bar{K}$ ) is nearly sufficient
- ▶  $R^2 > 0.99999$  in original KS paper
- ▶ Intuition:
  - ▶ Employed agents (high income) save similarly
  - ▶ Unemployed agents (low income) dissave similarly
  - ▶ Individual differences average out in aggregation
  - ▶ Distribution shape matters little for aggregate  $K'$
- ▶ This is called approximate aggregation
- ▶ Does NOT always hold (e.g., with larger shocks or more nonlinearity)